

# Public Cloud for Bioinformatics

Bioinformatics Workshop  
April 25th, 2022

# Data Sharing Needs for Biomedical Research

- Downloading large scale data to a local computer for analysis may be more difficult than bringing computational capability to where the data is located.
  - 1000 Genomes (>3,000, 80TB, > 2 days)
  - TopMed (>50,000, 2PB, weeks or months)
- Datasets are outgrowing local infrastructure, inhibiting researchers' ability to maintain them.
  - CCDG ~140,000 Whole Genome Sequence (WGS)
- Compute requirements to process these large datasets are exceeding local capacity, inhibiting analysis.
- Collaborating on research projects across organizations can be challenging due to differences in local IT environments.
  - Lack of standards and provenance
  - Batch effects, incompatible or inaccurate software
- Siloed and redundant infrastructure required for each lab, department, institution, etc.
  - Data Centers, Systems Administrators, Software Installation, User Support

# Public Cloud Options

## **Direct Cloud Usage (IaaS)**

*(more features, more control, harder to use/manage)*

Major Players:

1. Amazon Web Services (AWS) *(market leader)*
2. Microsoft Azure
3. Google Cloud (GCP: aka “Google Cloud Platform”)

## **Indirect Cloud Usage (PaaS or SaaS)**

*(less control, application focused, easier to use/manage)*

Examples:

- [Terra](#) (scientific computing; A joint Broad/Google Venture; see later discussion)
- [DNAnexus](#) (bioinformatics cloud platform; based on AWS)
- [Seven Bridges Genomics](#)
- [Google Colab](#) (scientific computing)
- [Heroku](#) / [Tableau](#) (websites & visualizations)
- [RIS](#) (they are working on “expanding” the cluster with the major cloud providers)

# Choosing a Cloud Option

It's very hard to be “cloud independent”. Inter-cloud operations are generally not encouraged by the cloud providers.

Much like choosing a programming language for your problem, choosing a cloud provider can impact your strategy, collaborations, and mindset.

Each provider has its own distinct features, tooling, and community.

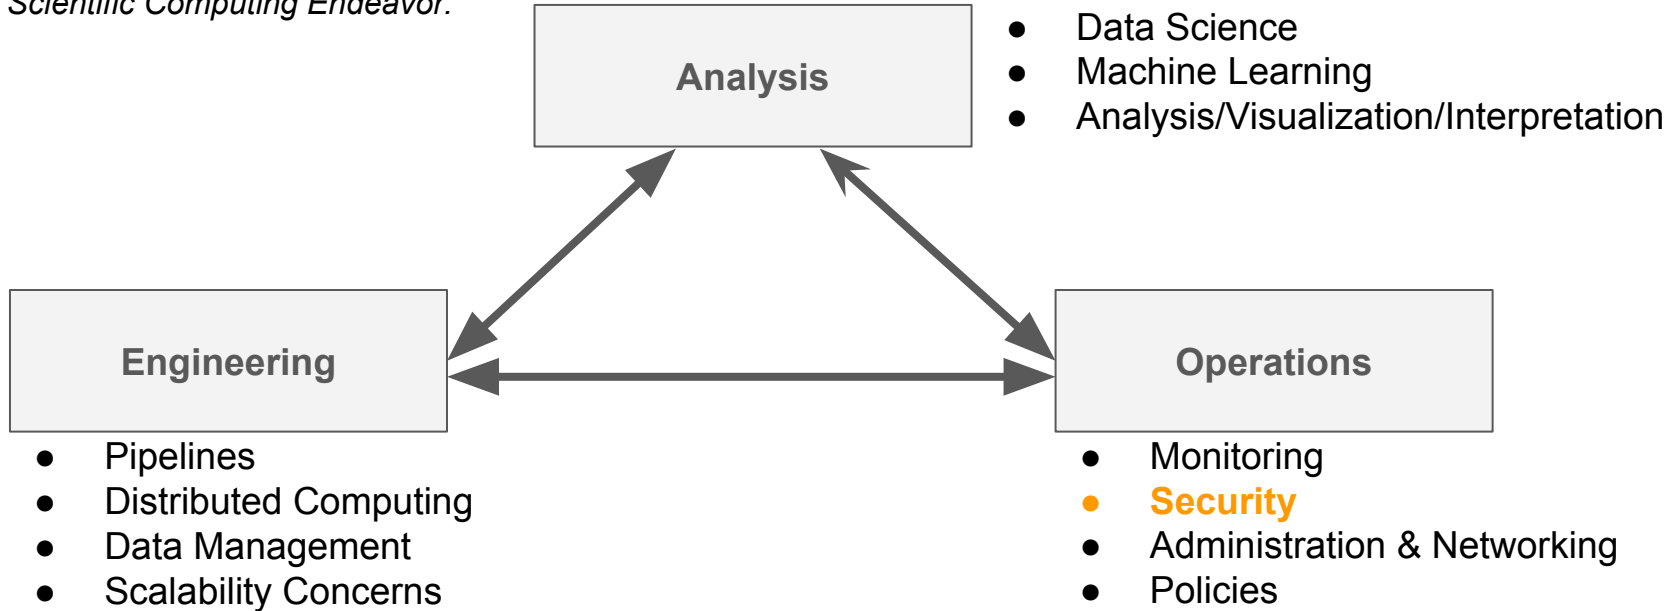
- Amazon
  - 80% market-share; very popular; more business-centric
- Microsoft Azure
  - enterprise-centric; WUIT preferred for protected data
- Google Cloud
  - expertise in Big Data
  - used by the Broad & [Data Biosphere](#) (bioinformatics tooling)
  - 3rd-place market-share; aggressively trying to increase it
  - WUIT-managed, dedicated fiber network to GCP for comparatively fast data transfers

# Public Cloud & in-house compute

- Using the cloud is not the same as using the compute0/1 cluster
  - Not HPC-oriented (no LSF, à la `bsub`; no NFS-based file system)
  - Essence of scientific cloud computing:
    - i. Create a VM
    - ii. Download Input Data onto VM from Cloud Storage
    - iii. Download software onto VM
    - iv. Run computation(s)
    - v. Upload Output Data on VM to Cloud Storage
    - vi. Delete VM
- “research” on the cloud vs. local in-house compute
  - Exploratory compute on the cloud can get expensive very quick.
  - Try to do small-scale “research” on in-house compute, and then move to the cloud for large-scale “production” tasks
- infrastructural considerations
  - see the next slide


# Infrastructural Considerations

*A Scientific Computing Endeavor:*



- Currently, on the cloud you are responsible for the full infrastructure
- Address either in-house or by purchasing 3rd party options -- in terms of personnel, tooling, security, and policy needs

# WUIT Public Cloud: Overview



- Request Cloud account from WUIT
- Request purchase order from department

- Test login, user permissions
- Access console features
- Request additional users, permissions from WUIT (as needed)

- Routing of bills
  - Email
  - Workday
- Billing Resources
- Billing Alerts
  - Daily
  - Monthly

- Minimize Cost
- Spot/Preemptible Compute
- Workflow Refactor
- Storage Footprint

# WUIT Public Cloud: Service Providers

- [Google Cloud - Information Technology](#)
  - [Request new WashU Google Cloud Project - Information Technology](#)



- [Amazon Web Services - Information Technology](#)
  - [Request New WashU AWS Account - Information Technology](#)



- [Microsoft Azure - Information Technology](#)
  - [Request New WashU Azure Subscription - Information Technology](#)





# WUIT Public Cloud: Initial Setup

- Required Information
  - School/Department/Unit Name
  - Project/Lab Name
  - Lead Technical Contact: Name, Email, Phone
  - Business/Finance Contact: Name, Email, Phone
  - Admin Access List: List of @wustl.edu addresses (additional to Lead Tech Contact)
  - Blanket Purchase Order with suppliers
- Vendor/Supplier Specific
  - AWS & Azure - HIPAA or FERPA protected data
  - AWS - Business-Level Support (10% additional fee)

# WUIT Public Cloud: Blanket Purchase Order (BO)

- Create through WorkDay/Marketplace
- The suppliers are:
  - Microsoft Azure - Insight Public Sector
  - Amazon Web Services - DLT SOLUTIONS LLC
  - Google Cloud - Burwood Group, Inc.
- Estimate your annual usage costs
  - Use a high estimate of what you expect to use
  - You decide the dollar amount, e.g. \$500.00
- Monthly invoices paid against the BO
  - Available dollars are correspondingly reduced
  - Change Order to a BO can be done to add money if necessary.
- BO's expire at the end of the fiscal year
  - Funds remaining on the BO go back into your department's budget.
  - There is no roll-over to the next fiscal year.
  - A new BO is issued for the new fiscal year.
- Set the expiration date as 6/30 of the fiscal year
- Description of the product as “[Cloud Provider] for [someone].”

# WUIT Public Cloud: Billing, POs, encumbering ...

- Billing contacts within each department open a Blanket Purchase Order (BO)
  - Billing contacts by Department
    - McDonnell Genome Institute
      - MGI\_purchasing <mgi\_purchasing@gowustl.onmicrosoft.com>
    - Medicine - may vary by PI
      - Ex. LaDonna Utley ([ladonna.utley@wustl.edu](mailto:ladonna.utley@wustl.edu))
- You may choose one BO per supplier or cloud/research project specific BOs
  - Check with Billing Contact for preference
- You will be asked periodically to encumber the BO for a certain amount based on how much you project to spend over the next 6-12 months
  - At this point you decide how much will be sourced to which fund numbers

# WUIT Public Cloud: Reseller Relationships

- For billing and legal purposes WUIT has established business agreements with resellers
  - DLT (AWS)
  - Burwood Group (GCP)
  - Insight Public Sector (Azure)
- Interaction with these resellers can be quite limited
  - Does add some complexity to billing because what you see in the cloud billing console is not necessarily what you will be charged by DLT, Inisght, or Burwood (may be less or more)

# WUIT Public Cloud: Cloud Computing Services Comparison

	WashU Private Cloud	Azure	AWS	Google Cloud
Signed University contract	✓	✓	✓	✓
University billing	✓	✓	✓	✓
WUSTL Key SSO login to admin console	N/A	✓	✓	✓
Approved for confidential and protected workloads	✓	*	*	*
WashU IT managed workloads	Servers, Databases, Web tools	Servers, Databases, Web tools		
Discounts	Competitive University pricing	6% across the board, Hybrid Use Benefit	5% across the board	25% on compute/storage, 19% on other services, Sustained use discounts
Data egress waiver	No egress data charges	No data egress charges if < 15% of total charges	No data egress charges if < 15% of total charges	No data egress charges if < 15% of total charges
Ease of integration with:	On-campus technology resources	Microsoft/Windows technology stack	Linux/Unix Technology stack	Linux/Unix Technology stack

\* WashU's agreements for Microsoft Azure, Amazon Web Services, and Google Cloud include HIPAA BAA and FERPA compliance guarantees. However, workloads (servers, databases, etc.) in these environments are not inherently compliant. Workloads must be built and configured following appropriate guidelines. When you will be using public cloud platforms to store and transmit protected data, It is necessary to seek guidance from Information Security and WashU IT.

If you still aren't certain which cloud offering would best suit your business needs, **contact** [<washuit-publiccloud@wustl.edu>](mailto:washuit-publiccloud@wustl.edu) to discuss the options.

Source:

<https://it.wustl.edu/services/cloud-computing/cloud-computing-services-comparison/> (4/21/2022)

# WUIT Public Cloud: Agreements

Normally, putting data into the cloud (ingress) is free. Taking data back out of the cloud (egress) is not typically free. If egress is less than 15% of charges, egress fees are waived.

WUIT has made agreements with the major public cloud providers to allow certain compute & storage discounts on WUIT-managed cloud accounts.

See WUIT's [Cloud Computing Service Comparison](#) for the latest details.

# Logging in to AWS (<http://connect.wustl.edu/awsconsole>)



Select a role:

▼ Account: 162241963241

☐ **Genome-AWS-MGriffithLab-Admin**

▼ Account: 182652044697

☐ **Oncology-AWS-GriffithOpenData-Admin**

▼ Account: 393523350052

☐ **WUSM-AWS-CottoTrainingAward-Admin**

▼ Account: 557711160097

☐ **Genome-AWS-OGriffithLab-Admin**

▼ Account: 653839114923

☐ **Genome-AWS-GriffithLab-Admin**

Sign in

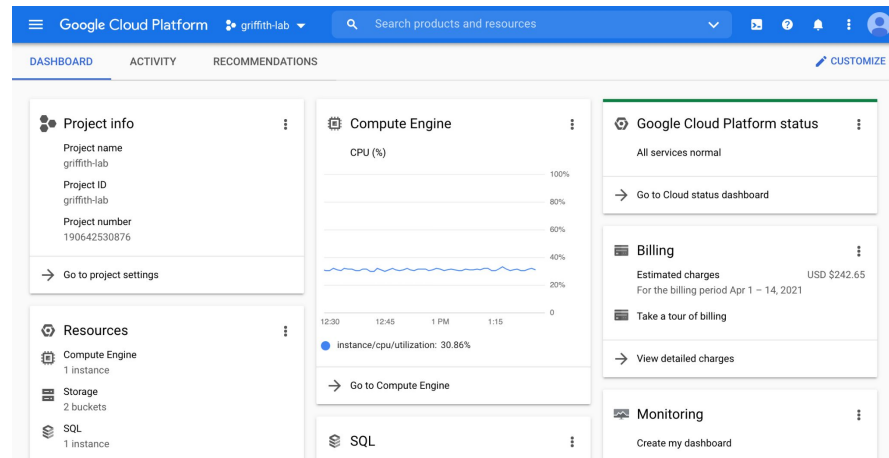
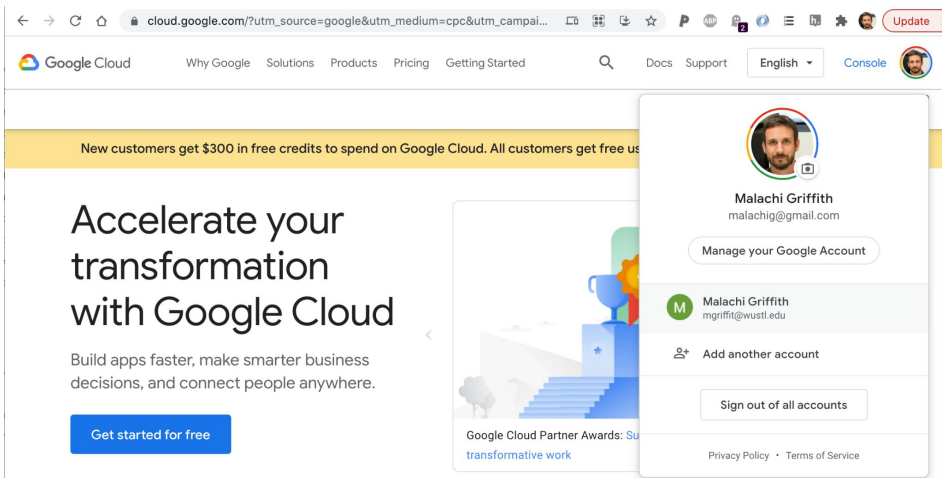
English ▼

[Terms of Use](#) [Privacy Policy](#) © 1996-2021, Amazon Web Services, Inc. or its affiliates.

- Authentication uses your WUSTL Key
- You can have multiple accounts (optional)
  - Per project
  - Per billing source
- Or you have can one account and manage billing by paying one bill from multiple sources
- Multiple accounts can also be used to compartmentalize access to members of the lab working on different projects
- Each account will generate distinct billing, billing alerts, etc.

# Logging into GCP (<https://cloud.google.com/>)

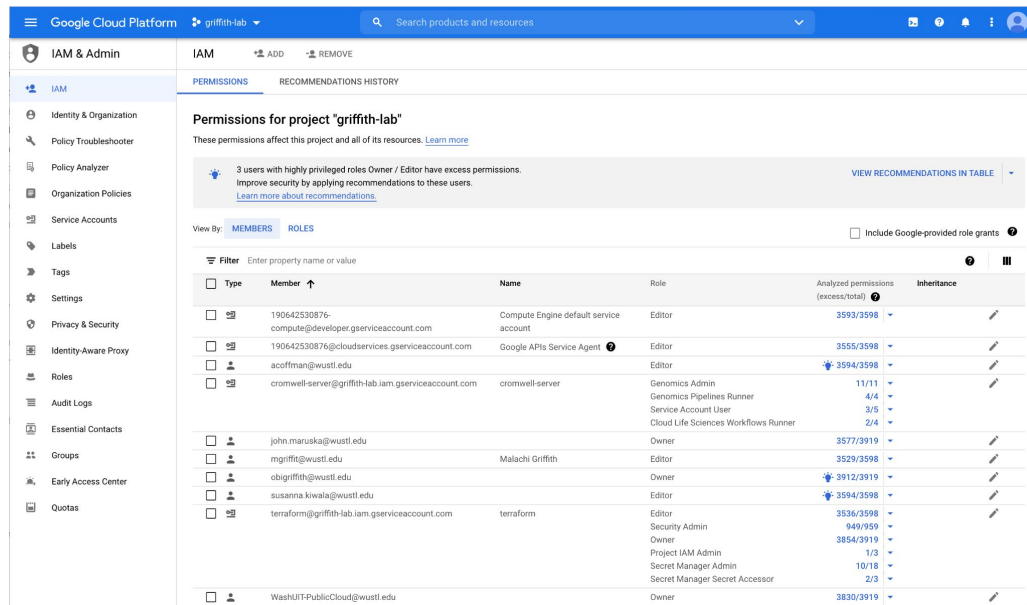
- WUSTL Key Google account
  - A very limited functionality Google account associated with your WASHU login
  - You now have two Google accounts. If you use the Google ecosystem a lot, this will be annoying.
  - Better functionality would happen with [Google Workspace \(formerly known as “G Suite”\)](#) integration; however, WUIT has provided G Suite access to some users, contact WUIT public cloud for info.





# Adding users from a lab to a cloud account

- For AWS, the ability to add users from the lab to access AWS accounts is managed through WUIT (Eric Suiter) by email
  - No way for PI to see what users are currently authorized without asking WUIT
- For GCP, the ability to add users is also limited to WUIT but at least it is possible to view the current configuration



The screenshot displays the Google Cloud Platform IAM & Admin console for the project 'griffith-lab'. The left sidebar shows the navigation menu with 'IAM' selected. The main content area is titled 'Permissions for project "griffith-lab"' and shows a list of members under the 'MEMBERS' tab. A notification banner indicates that 3 users with highly privileged roles (Owner / Editor) have excess permissions, suggesting a security improvement by applying recommendations.

Type	Member	Name	Role	Analyzed permissions (excess/total)	Inheritance
Service Account	190642530876-compute@developer.gserviceaccount.com	Compute Engine default service account	Editor	3593/3598	
Service Account	190642530876@cloudservices.gserviceaccount.com	Google APIs Service Agent	Editor	3555/3598	
User	acoffman@wustl.edu		Editor	3594/3598	
Service Account	cromwell-server@griffith-lab.iam.gserviceaccount.com	cromwell-server	Genomics Admin Genomics Pipelines Runner Service Account User Cloud Life Sciences Workflows Runner	11/11 4/4 3/5 2/4	
User	john.maruska@wustl.edu		Owner	3577/3919	
User	mgriffith@wustl.edu	Malachi Griffith	Editor	3529/3598	
User	obgriffith@wustl.edu		Owner	3912/3919	
User	susanna.kiwala@wustl.edu		Editor	3594/3598	
Service Account	terraform@griffith-lab.iam.gserviceaccount.com	terraform	Editor Security Admin Owner Project IAM Admin Secret Manager Admin Secret Manager Secret Accessor	3536/3598 949/959 3854/3919 1/3 10/18 2/3	
User	WashJUIT-PublicCloud@wustl.edu		Owner	3830/3919	

# Invoice example



2411 Dulles Corner Park  
Suite 800  
Herndon, VA 20171

Invoice Questions: 888-358-9346  
General Information: 703-709-7172  
Fax: 866-352-5855

Tax ID No: 54-1599882  
CA Reseller: 101643630  
DB No: 78-6468199

Invoice No.: SI511473  
Order: 4900439  
Customer: WUN10  
Contract #: N/A

GST No: 82690 0003 RT0001  
MB PST No: 826900003MT0001  
SK PST No: 2476547  
QST No: 1217287088

Invoices will be emailed to you monthly for approval. Work with DLT to ensure the right people get emailed.

**Bill To:** Washington University in St. Louis  
Accounts Payable  
Campus Box 1056  
700 Rosedale Ave  
SAINT LOUIS, MO 63112-1408

**Ship To:** Washington University in St. Louis  
Accounts Payable  
ONCOLOGY / SCB  
425 SOUTH EUCLID  
SOUTHWEST TOWER  
SAINT LOUIS, MO 63108

- Performance Period
- WASHU PO
- AWS Account #

Date		Period of Performance		ACT # / PDN #		Terms	
02/26/21		01/01/21 - 01/31/21				Net 30 Days	
Purchase Order Number			Order Date	Salesperson		Our Order Number	
2940074K			02/26/21	Jennifer Triplett		4900439	
Quantity		CLIN No.	Item Number		Tax	Unit Price	Amount
Req.	Ship	B.O.	Description				
1	1		413500		N	1,491.27	1,491.27

AWS 01-2021 - WUSTL McDonnell Genome Institute

# Invoice example

Description	Qty	Rate	Price
4900439 - 6538-3911-4923 - McDonnell Genome Institute - Genome Griffith Lab			
<b>CloudWatch</b>			
Alarm metric month (standard resolution) - US West (Oregon)	2	0.09500	0.19
GB-mo of log storage - US East (Northern Virginia)	0.01	0.00000	-
GB-mo of log storage - US West (Oregon)	0	0.00000	-
<b>CloudWatch Total</b>	<b>2.01</b>	<b>0.03167</b>	<b>0.19</b>
<b>DynamoDB</b>			
GB - first 10 TB / month data transfer out beyond the global free tier	0	0.00000	-
GB - next 40 TB / month data transfer out	0	0.00000	-
<b>DynamoDB Total</b>	<b>0</b>	<b>0.00000</b>	<b>-</b>
<b>EC2</b>			
Elastic IP address not attached to a running instance per hour (prorated)	2231	0.00475	10.60
EUN1-AWS-Out-Bytes in EU (Stockholm)	0	0.00000	-
GB - first 10 TB / month data transfer out beyond the global free tier	1738.81	0.08550	148.67
GB - regional data transfer - in/out/between EC2 AZs or using elastic IPs or ELB	26.48	0.00944	0.25
GB - US West (Oregon) data transfer to Asia Pacific (Mumbai)	0	0.00000	-
GB - US West (Oregon) data transfer to Asia Pacific (Seoul)	0.01	0.00000	-
GB - US West (Oregon) data transfer to Asia Pacific (Singapore)	0.23	0.00000	-
GB - US West (Oregon) data transfer to Asia Pacific (Sydney)	0	0.00000	-
GB - US West (Oregon) data transfer to Asia Pacific (Tokyo)	0.06	0.00000	-
GB - US West (Oregon) data transfer to AWS GovCloud (US)	0.02	0.00000	-
GB - US West (Oregon) data transfer to Canada (Central)	0.23	0.00000	-
GB - US West (Oregon) data transfer to EU (Germany)	0.01	0.00000	-
GB - US West (Oregon) data transfer to EU (Ireland)	1.3	0.01538	0.02
GB - US West (Oregon) data transfer to EU (London)	0	0.00000	-
GB - US West (Oregon) data transfer to EU (Paris)	0	0.00000	-
GB - US West (Oregon) data transfer to South America (Sao Paulo)	0	0.00000	-
GB - US West (Oregon) data transfer to US East (Northern Virginia)	19.4	0.01907	0.37
GB - US West (Oregon) data transfer to US East (Ohio)	0.27	0.03704	0.01
GB - US West (Oregon) data transfer to US West (Northern California)	0	0.00000	-
GB-month of General Purpose SSD (gp2) provisioned storage - US West (Oregon)	2560	0.09500	243.20
GB-Month of snapshot data stored - US West (Oregon)	1824.65	0.04750	86.67
In Asia Pacific (Hong Kong)	0.02	0.00000	-
In Middle East (Bahrain)	0	0.00000	-
On Demand Linux m5.2xlarge Instance Hour	744	0.36480	271.41
On Demand Linux m5a.2xlarge Instance Hour	744	0.32680	243.14
On Demand Linux t2.large Instance Hour	744	0.08816	65.59
On Demand Linux t2.medium Instance Hour	744	0.04409	32.80

- DLT gives a detailed breakdown of cloud services you are being charged for
- This only loosely corresponds to the billing information you see in AWS console
  - The usage should be the same, but \*how\* you are billed for it can be very different
  - We have seen DLT bills >2x what is shown in AWS
  - Related to reserved instances. DLT calls this “RI arbitrage”. Basically DLT is reserving instances, keeping the savings from Amazon and billing you the full agreed rate.

# CloudCheckr CMx (<https://app-us.cloudcheckr.com/>)

- More detailed information on what DLT will actually charge you for AWS usage
- Must request that DLT set this up for you

## CloudCheckr CMx



6538391... AWS

### Billing & Cost Management Dashboard


Select Month: Apr 2021 ☒ Include RI Upfront Costs

Filter


Current Month-to-date for April 2021

**\$379.61**

#### Spend Summary



#### Month-to-Date Spend by Service



CloudCheckr  
Console

**Getting Started with AWS Billing & Cost Management**

- Manage your costs and usage using [AWS Budgets](#)
- Visualize your cost drivers and usage trends via [Cost Explorer](#)
- Dive deeper into your costs using the [Cost and Usage Reports](#) with [Athena integration](#)
- **Learn more:** Check out the [AWS What's New webpage](#)

**Do you have Reserved Instances (RIs)?**

- Access the [RI Utilization & Coverage reports](#)—and [RI purchase recommendations](#)—via [Cost Explorer](#).


#### Spend Summary

Cost Explorer

Welcome to the AWS Billing & Cost Management console. Your last month, month-to-date, and month-end forecasted costs appear below.

Current month-to-date balance for April 2021

**\$274.29**



#### Month-to-Date Spend by Service

Bill Details

The chart below shows the proportion of costs spent for each service you use.

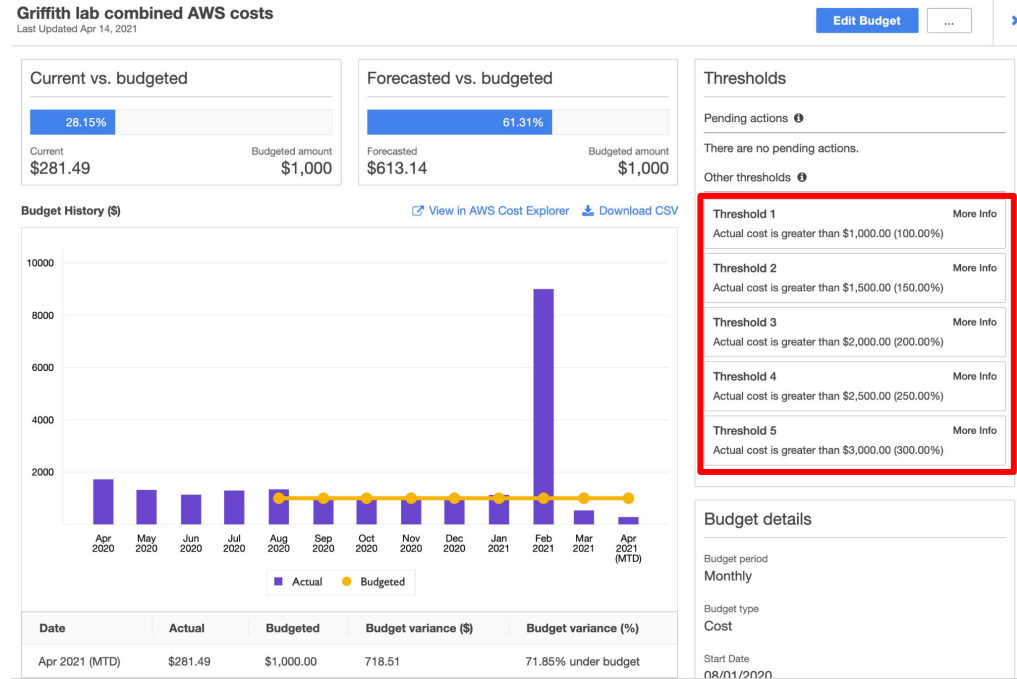


EC2	\$199.24
DataTransfer	\$74.77
RDS	\$0.19
CloudWatch	\$0.09
Other Services	\$0.00
Tax	\$0.00
Total	\$274.29

AWS Billing  
Console

# “Budgets” and Billing Alerts

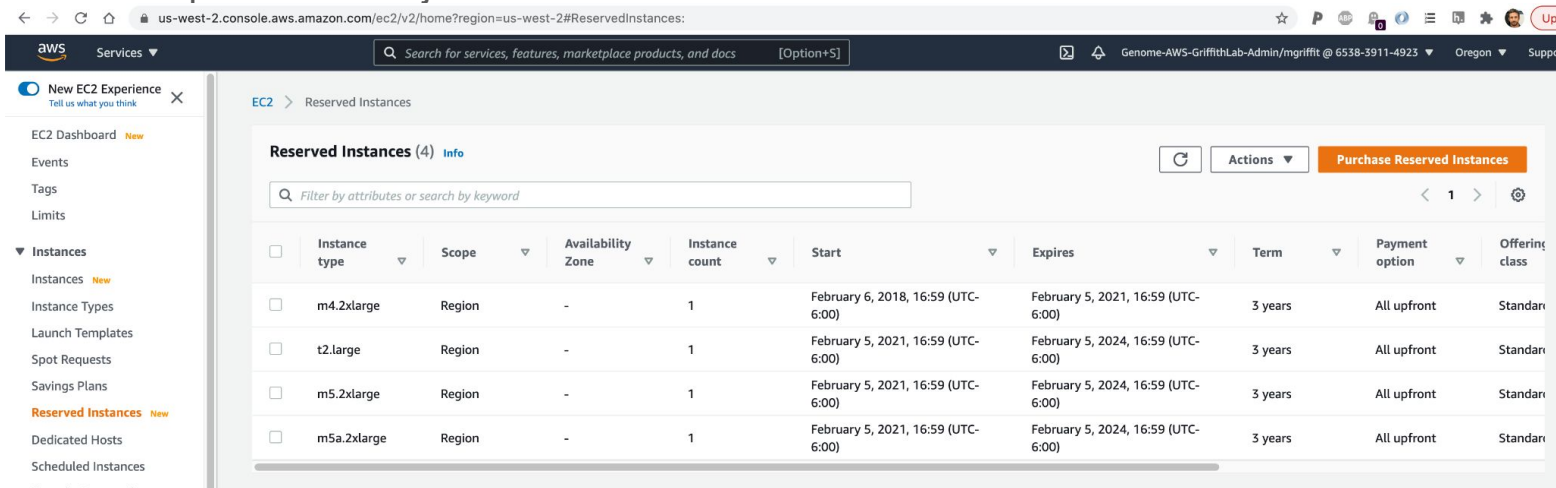
- Creating a “budget” in AWS or GCP allows you to set up email/text alerts when you hit certain amounts each month



# Long-Term Compute:

## “Reserved Instances” (AWS) and “Commitments” (GCP)

- If you have persistent compute services running use of Reserved Instances or Commitments can result in huge savings
  - E.g. by reserving instances for 3 years at a cost of ~\$8,000 we project savings of ~\$15,000 compared to if we just used on demand resources for that whole time



The screenshot shows the AWS Management Console interface for Reserved Instances. The breadcrumb navigation indicates the path: EC2 > Reserved Instances. The main heading is "Reserved Instances (4) Info". Below this is a search bar with the placeholder text "Filter by attributes or search by keyword". To the right of the search bar are buttons for "Actions" and "Purchase Reserved Instances". Below the search bar is a table listing four reserved instances.

	Instance type	Scope	Availability Zone	Instance count	Start	Expires	Term	Payment option	Offering class
<input type="checkbox"/>	m4.2xlarge	Region	-	1	February 6, 2018, 16:59 (UTC-6:00)	February 5, 2021, 16:59 (UTC-6:00)	3 years	All upfront	Standard
<input type="checkbox"/>	t2.large	Region	-	1	February 5, 2021, 16:59 (UTC-6:00)	February 5, 2024, 16:59 (UTC-6:00)	3 years	All upfront	Standard
<input type="checkbox"/>	m5.2xlarge	Region	-	1	February 5, 2021, 16:59 (UTC-6:00)	February 5, 2024, 16:59 (UTC-6:00)	3 years	All upfront	Standard
<input type="checkbox"/>	m5a.2xlarge	Region	-	1	February 5, 2021, 16:59 (UTC-6:00)	February 5, 2024, 16:59 (UTC-6:00)	3 years	All upfront	Standard

[RIS CONFLUENCE LINK \(RESTRICTED ACCESS\)](#)


# Short-Term Compute:

## “Spot” (AWS) or “Preemptible” (GCP) Instances

- Meant for “short-term” fault-tolerant jobs
- Much cheaper pricing (up to 90% at times) than regular compute instances
- Drawback is that AWS or GCP reserves the right to take back the machine at any given time depending on the state of the cloud
- Google Preemptible Instances can run at most for 24 hours
- See [GCP Preemptible VM Instances](#) or [Amazon EC2 Spot](#) for more details
- Often useful for running many batch jobs at scale cheaply

# Public Cloud Resources: Education & Training

- [Introduction to AWS | Griffith Lab](#)
- [Getting Started with AnVIL](#)
- [Training | NIH STRIDES](#)
- [Google Cloud Online Courses | Coursera](#)
- [Learn AWS with Training and Certification | Cloud Skills Courses and Programs](#)
- [Genomics in the Cloud \[Book\]](#)






# Public Cloud Resources: Funding Opportunities

- Current
  - [Enrollment | NIH STRIDES](#)
  - [AWS Programs for Research and Education](#)
  - [Get and redeem education credits | Google Cloud Billing](#)
- Past
  - [AnVIL Cloud Credits Continued Program \(AC3\)](#) (12/17/2021)

# Cloud Work Stream Vision



Global Alliance  
for Genomics & Health



# Cloud Work Stream APIs



Global Alliance  
for Genomics & Health

**Sharing Tools  
and Workflows**


**Executing  
Workflows**

**Executing  
Individual  
Tasks**


**Accessing Data**




Tool Registry Service (TRS)




Workflow Execution Service (WES)



Task Execution Service (TES)



Data Repository Service (DRS)




# Tool Registry Service (TRS) API



Global Alliance  
for Genomics & Health


*List, search, & config CWL/WDL/Nextflow/Galaxy-described Docker Tools and Workflows.*

CWL/WDL/  
Nextflow/Galaxy  
Workflow



&

CWL/WDL/Nextflow/  
Galaxy-Described Tools



TRS Sharing API

GET list

GET search

GET config

**Dockstore,  
Biocontainers,  
& other repos**

- **GitHub page:** <https://github.com/ga4gh/tool-registry-service-schemas>
- **Latest release:** 2.0.1-beta.0
- **Champion:** Denis Yuen

**Official GA4GH  
Standard!**

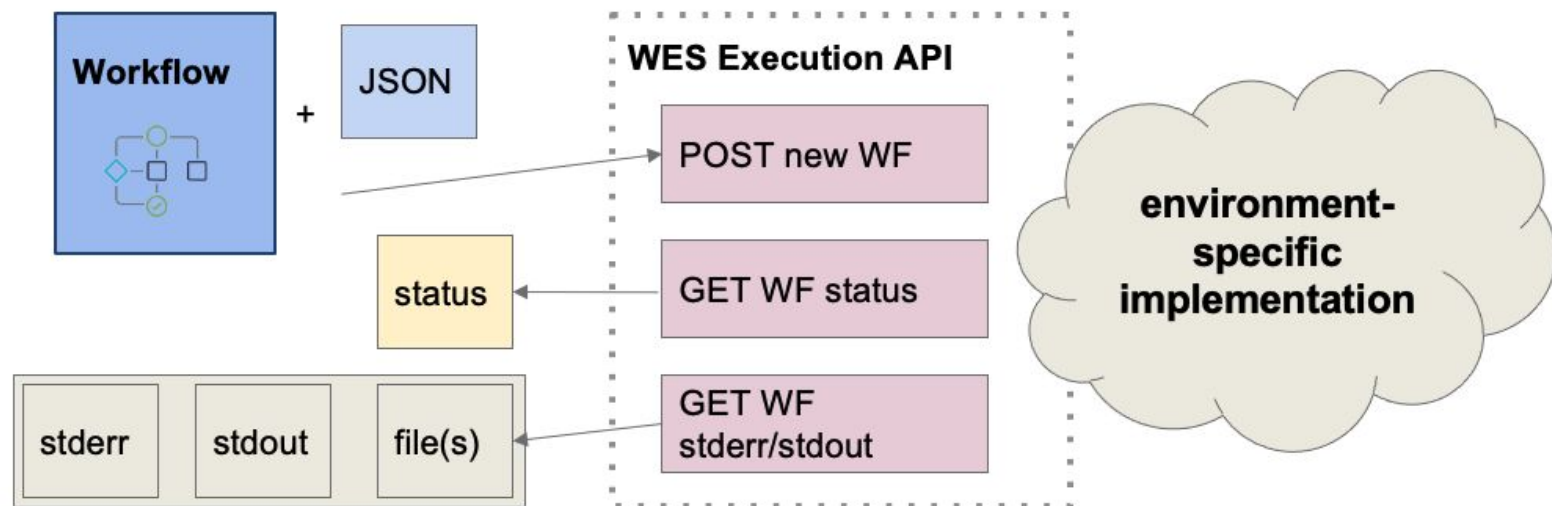
ga4gh.org

# Workflow Execution Service (WES) API



Global Alliance  
for Genomics & Health

Execute Workflows in a cloud and platform-agnostic way. *(TES is very similar!)*



- **GitHub page:** <https://github.com/ga4gh/workflow-execution-service-schemas>
- **Latest release:** 1.0.0
- **Champions:** James Eddy, Ruchi Munshi (former), and Walt Shands

Official GA4GH  
Standard!

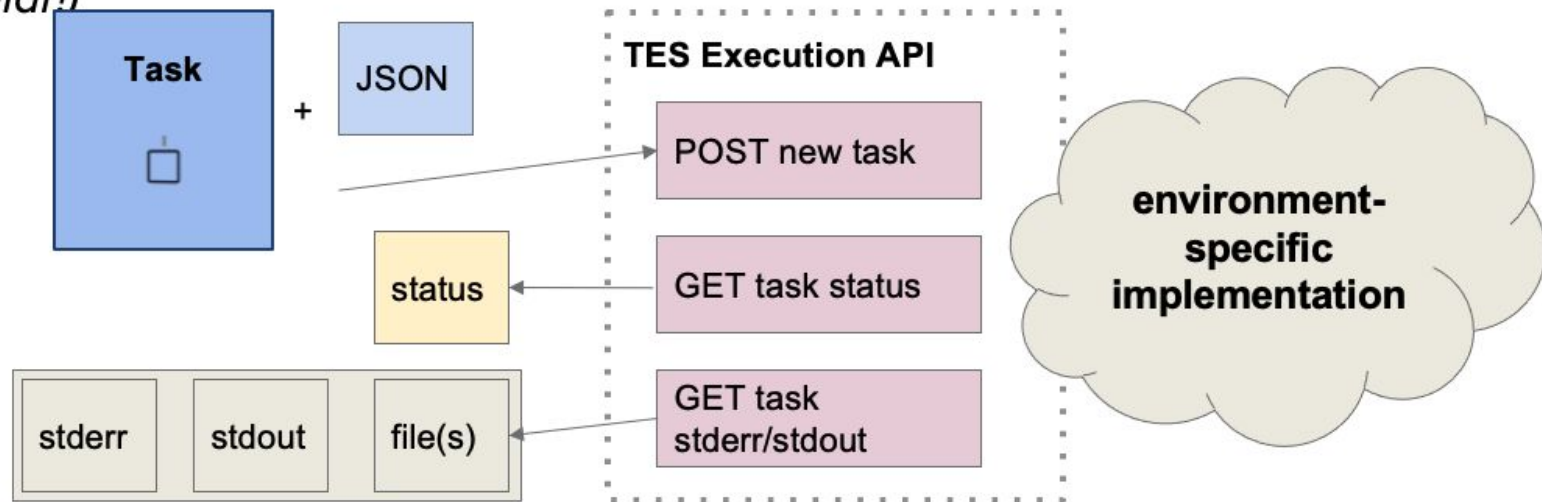
...org

# Task Execution Service (TES) API



Global Alliance  
for Genomics & Health

Execute individual Tasks in a cloud and platform-agnostic way. (*WES is very similar!*)



- **GitHub page:** <https://github.com/ga4gh/wiki/wiki/Task-Execution-Service>
- **Latest release:** 1.0.0
- **Champions:** Ania Niewielska and Kyle Ellrott

Official GA4GH  
Standard!

...org

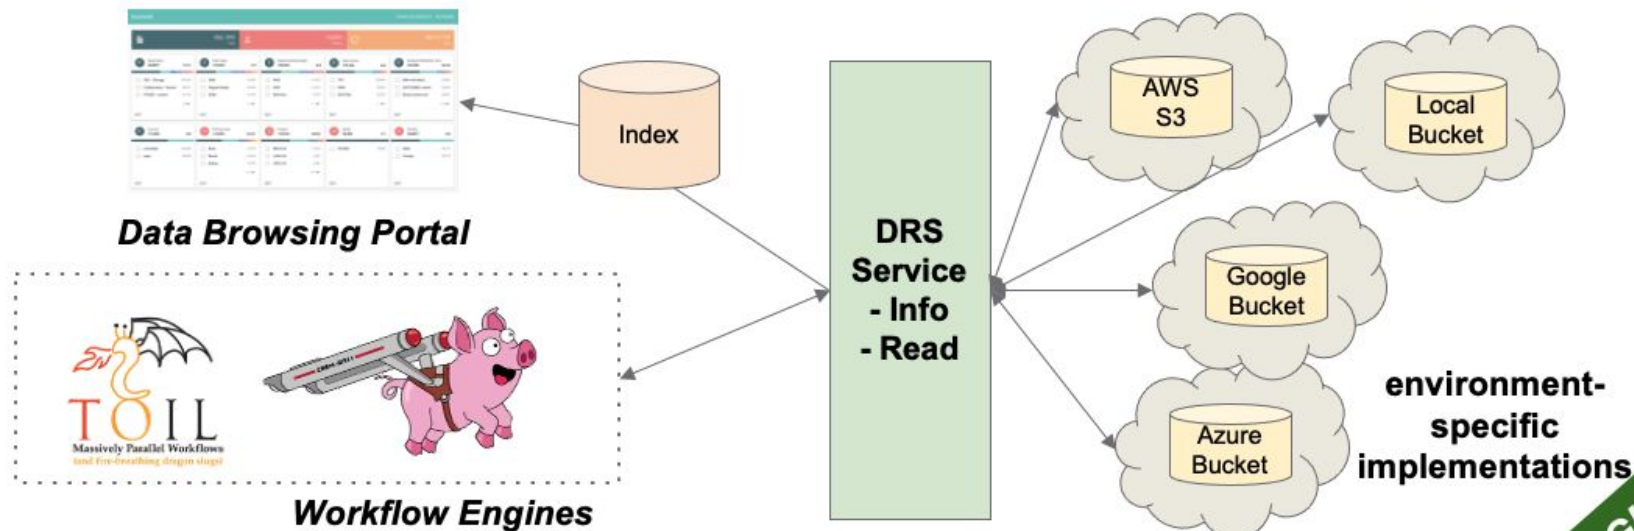


# Data Repository Service (DRS) API



Global Alliance  
for Genomics & Health

*A cloud agnostic way to map IDs to data objects and get info/read via signed URLs or other protocols*



- **GitHub page:** <https://github.com/ga4gh/data-repository-service-schemas>
- **Latest release:** 1.2.0
- **Champions:** David Glazer and Brian O'Connor

**Official GA4GH  
Standard!**

ga4gh.org


# Analysis Visualization and Informatics Lab-space (AnVIL)

Inverting the model of genomics data sharing with the NHGRI Analysis Visualization and Informatics Lab-space (AnVIL)




*AnVIL empowers users to collaboratively analyze data in a secure cloud environment*

## Data



Large, Diverse  
Datasets

## Compute

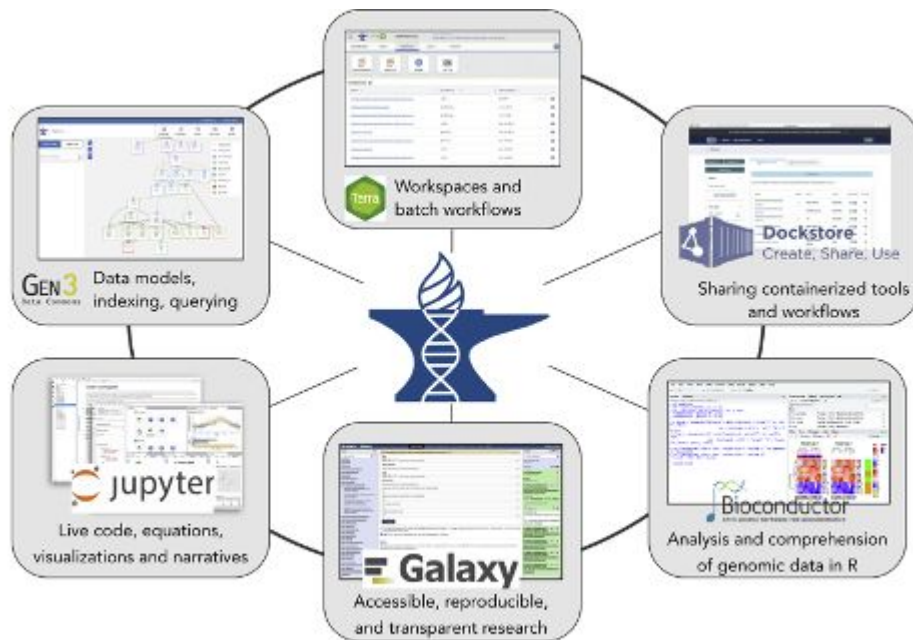


Secure &  
Reproducible

## Users



Democratized &  
Collaborative



Source: "Inverting the model of genomics data sharing with the NHGRI Genomic Data Science Analysis, Visualization, and Informatics Lab-space":  
<https://doi.org/10.1016/j.xgen.2021.100085>



# NIH Cloud Platform Interoperability (NCPI) Effort



## RAS: Researcher Auth Service

Goal: Unified identity/authentication

RAS is an effort by the NIH's Center for Information Technology (CIT) to provide a common mechanism by which researchers can establish their identity and access data they are authorized to use.



## DRS: Data Repository Service

Goal: Unified data access across storage infrastructures

[Data Repository Service API](#), are a standardized set of cloud data access methods. The primary functionality is to map a logical ID to a means for physically retrieving the data represented by the `drs://URI` scheme.



## FHIR: Fast Healthcare Interoperability Resources

Goal: data harmonization and an API for exchange of electronic medical records.

FHIR facilitates interoperation between health care systems, to make it easy to provide health care information on a wide variety of devices.

